AI for IPCC:
Towards Living Evidence Synthesis for Climate Science

Sandeep Chowdhary[1,*]

[1]International Institute for Applied Systems Analysis, Laxenburg, Austria.
[*]chowdhary@iiasa.ac.at

ABSTRACT

Scientific synthesis is facing a bottleneck: the volume of literature grows exponentially while expert attention remains finite. Global assessments such as the IPCC exemplify this challenge, with multi-year lags between new findings and their incorporation into policy-relevant reports; a gap that will only widen if publication rates continue to accelerate. We propose a transparent, open-source, *living* evidence system that continuously retrieves, verifies, and integrates climate findings into a versioned HTML IPCC report, preserving sentence-level source and validator scores. The pipeline combines citation-graph expansion, semantic retrieval, and LLM-based relevance classification; applies dual-model claim extraction with cross-validation; and escalates low-confidence outputs to human reviewers. We will benchmark performance along three axes: (i) coverage gains at fixed precision relative to curated IPCC reference lists, (ii) end-to-end accuracy on expert-labelled test sets (precision, recall, calibration, validator stability), and (iii) decision utility measured as time-to-evidence and source trust, compared to both leading AI tools and manual curation. Horizon-scanning modules (burst detection, novelty tracking, and atypical-combination signals) will identify emerging themes and blind spots post-AR6, while an API will link validated findings to Integrated Assessment Models. By replacing episodic manual synthesis with auditable continuous integration, the system aims to lower screening costs, shorten evidence-to-policy latency, and generalise to other evidence-intensive domains.

1 State of the Art

Climate change stands among the most profound challenges of our time. Since the first quantitative prediction of CO2-induced warming [1] and early estimates of climate sensitivity [2, 3], research has expanded to encompass economic impacts [4], carbon budgets [5], and planetary boundaries [6], reaffirming both the severity of the crisis and the urgency of global responses.

Global environmental assessments (GEAs) such as the IPCC, IPBES, GEO, UNEP's Gap reports, and the Global Sustainable Development Report provide crucial evidence on environmental change and guide action-oriented policy [7]. They have shaped international negotiations and agreements, including the Paris Agreement [8, 9]. The IPCC in particular has long informed climate governance, with AR6 highlighting the urgency of mitigation and adaptation and feeding into the Global Stocktake [10]. Yet concerns remain about their ability to keep pace with the rapidly expanding scientific literature [11, 7]. The number of publications grows exponentially doubling roughly every 15 years while human attention and processing capacity remain bounded [12, 13]. This creates an unavoidable tension: ever more knowledge is produced, but the ability of experts to keep pace with emerging findings is limited, producing a synthesis bottleneck.

The IPCC exemplifies both the strengths and limits of *manual synthesis*. Each cycle mobilises hundreds of experts over 5–7 years [14], during which tens of thousands of publications must be screened as the literature grows rapidly. In AR5, authors assessed about 30,000 publications [15], while in AR6 this doubled to more than 66,000 [16]. Working Group I in AR6 engaged 234 authors to review 14,000 papers and address 78,000 review comments [17], while Working Group III involved 278 lead authors and 354 contributing authors to assess 18,000 papers and respond to 59,000 comments [18]. Beyond this scale, assessments have expanded from physical science to impacts, mitigation, adaptation, and social dimensions [19, 20], drawing on disciplines from economics and psychology to engineering and the humanities, and reaching audiences from negotiators to NGOs and businesses worldwide [20]. The direct IPCC Trust Fund expenditure for 2024 exceeded €5 million [21], excluding in-kind contributions. Assessment writing is thus one of the most resource-intensive processes in global science, and the lag of several years between research and synthesis creates a bottleneck.

Recent advances in machine learning (ML) and natural language processing (NLP) promise to ease this bottleneck. Machine learning evidence maps of climate impacts [22] aggregated over 100,000 impact studies and revealed geographic attribution gaps, but they remain static and stop short of producing assessment-ready synthesis. AI-driven mapping of adaptation policies

[23] identified systematic differences across governance levels, yet highlighted persistent blind spots in vulnerable countries and lacked iterative updating. A "living" ML map of 84,000 mitigation policy studies [24] demonstrated disparities between research attention and emissions, but still lacked integration into modelling frameworks. AI-enhanced systematic mapping of carbon dioxide removal (CDR) literature [25] uncovered three times more relevant publications than prior estimates, but outputs were static and disconnected from IAM parameters. Large-scale bibliometric studies [26, 27] map research landscapes, but remain descriptive without mechanisms for continuous integration. Domain-specific models such as ClimateBERT [28] have improved classification in corporate climate disclosures, and transformers have enabled attribution mapping of climate impacts [22] and systematic reviews in adaptation and health [29, 30, 31]. These approaches also open pathways for multilingual evidence synthesis, addressing English-language bias through translation and cross-lingual classification [32, 33]. Living evidence synthesis platforms [34] and AI-supported screening [35] show promise, while frameworks for structured confidence evaluation [36, 37, 38] mirror and extend expert judgment in assessment processes. Yet problems remain: hallucination risks, frozen corpora that become outdated, and uneven performance in long-form synthesis [39, 40]. General-purpose chatbots (e.g., ChatGPT, Claude Sonnet, Gemini) provide fast answers but lack coverage, transparency, and reliability in long-form scientific tasks [41, 42]. Specialized AI-for-science tools, by contrast, target literature discovery and synthesis but face other limitations: most are proprietary and closed-source (see Table 1), with restricted free tiers, opaque methods, and limited synthesis scale. Few explicitly support consensus-building, underscoring the difficulty of achieving open, transparent, and assessment-ready AI-powered scientific discovery.

Therefore, what is missing is an approach that combines the breadth and scalability of automated evidence mapping with the precision and accountability required for policy ready assessments. What is needed is an open-source tool that can simultaneously expand coverage across the rapidly growing literature, anchor findings directly to their sources, and integrate them into the IPCC framework without introducing hallucinations or relying on black boxes. Only such a design can close the growing gap of literature synthesis improving timeliness while maintaining rigour in climate science.

Tool	Open Source	Free	Data Source/Coverage	Black Box	Synthesis Size	Summary	Targeted Consensus
Semantic Scholar[43]	No	Yes	Semantic Scholar corpus (200M+ papers, multidisciplinary)	No (API+UI)	Discovery only	Yes (TLDR)	No
Elicit[44]	No	Yes*, Paid	100M+ (Semantic Scholar, abstracts/ fulltext if open)	Mostly	10/25/40	Yes	Partial (evidence table)
Consensus[45]	No	Yes*, Paid	200M+ (Semantic Scholar, OpenAlex, web crawl)	Yes	1,500 filtered, top 20 synthesized	Yes	Yes (explicit claim meter)
Paperguide[46]	No	Yes*, Paid	200M+ (broad, not fully publicized)	Yes	~100	Yes	Yes (Q&A/synthesis)
Scite[47]	No	Free trial, Paid	Mixed (Elsevier, PubMed, others, open)	Part	Per-paper/field	No	No
Connected Papers[48]	No	Yes*, Paid	Semantic Scholar, ODC-BY, ResearchGate, Academia	Yes	~50,000 (graph)	No	No
Research Rabbit[49]	No	Yes (free)	OpenAlex, Semantic Scholar	Yes	Thousands, visual	No	No
Scholarcy[50]	No	Yes*, Paid	Uploaded papers only	Yes	Tens–hundreds	Yes	No
SciSpace[51]	No	Yes*, Paid	Multisource, open academic papers/metadata	Yes	5–20 (per synthesis)	Yes	No
Iris.ai[52]	No	Free*, Paid	Open, paywalled, enterprise data	Yes	Hundreds	Yes	No

Table 1. Feature comparison of top AI-for-science tools. (*=very limited)

2 Objectives

To address the synthesis bottleneck outlined above, this project will build an automated, systematic literature synthesis tool for climate science that keeps IPCC reports continuously updated with minimal manual supervision.

Can we achieve reliable, continuous literature integration into IPCC reports? We will design a hybrid human-AI workflow that automatically retrieves, validates, and inserts new climate science findings into a live HTML version of the IPCC reports, with source preserved down to sentence-level PDF spans.

How well does the system perform compared to manual curation and existing AI tools, and what does this mean for policy relevance and trust? We will benchmark performance against the gold standard of manual annotation in IPCC processes, testing retrieval, classification, and synthesis accuracy at each step of the pipeline. This includes (i) coverage gains

over curated IPCC reference lists at fixed precision, (ii) end-to-end accuracy via expert labelled test sets with precision/recall, calibration, and validator stability, and (iii) closeness of automated key takeaway insertion to expert authored IPCC text. In parallel, we will compare our system's outputs to existing ML-based bibliometric tools and general purpose LLMs (e.g., ChatGPT, Gemini, Claude) to establish relative strengths and weaknesses. The results will be assessed not only for technical accuracy but also for decision utility: time-to-evidence, trust in source cues, and perceived policy relevance by IPCC authors and modellers.

Can AI be used for horizon scanning for IPCC gaps and detect emerging blind spots?   AR6 highlighted clear priorities: in mitigation, more robust MRV for carbon dioxide removal, demand and services side options, and transparent scenario curation linked to national pathways [53, 54, 55, 56, 57, 58, 59]; in adaptation, stronger evidence on limits, cascading risks, and urban systems [60, 61]; and in physical science, better detection/attribution of extremes and compound events [62]. Our tool will operationalise these priorities by continuously tracking growth signals (e.g., citation bursts, novelty terms, and publication activity), quantifying post AR6 attention, and routing results into curated watchlists and "trend cards" with validator scores, timestamps, and links to source PDFs. Beyond reinforcing known priorities, it will also detect emerging blind spots: identifying novel lexicon in titles and abstracts, as well as atypical combinations of research domains. Using automated classification into existing IPCC themes, the system will surface not only extensions of current categories but also entirely new re-combinations patterns shown to drive high impact science [63].

3 Originality and Innovation

The central originality of this project lies in *combining multiple modes of search and reasoning into a single pipeline*. Specifically, we will integrate *systematic algorithmic discovery through citation networks*, *semantic similarity search*, and *LLM-based contextaware classification* to evaluate whether a new study is relevant to the *specific research topic, subsection, or hypothesis under review*. This layered approach ensures that directly related, indirectly connected, and contextually relevant literature are all captured, far beyond what keyword searches or one-off queries can achieve. By merging *graph expansion*, *semantic enrichment*, and *contextual decision-making*, the tool introduces a fundamentally new way to achieve comprehensive and precise synthesis.

Equally original is the emphasis on *transparency*. Current AI-supported literature tools are largely closed-source, proprietary, and opaque in how they retrieve and filter information. Systems like Illicit provide structured outputs but restrict corpus size, limit explainability, and prevent users from knowing whether key studies were omitted. Our approach explicitly inverts this model: rather than a *black box*, it is designed as a *glass box* where all steps from search boundaries to intermediate filtering decisions to sentence-level are open, audit-able, and user-configurable. This transparency not only builds trust but also enables replication and extension in other scientific domains.

Finally, the project goes beyond static reports and creates a *living knowledge system*. For centuries, books and bound reports were the primary vessels of scientific synthesis- fixed snapshots in time, necessarily outdated soon after they are published. The future, however, lies in *living, breathing documents*: continuously updated, automatically-validated, and openly accessible, with every claim transparently linked to its source. Our system will instantiate this future for climate science, integrating directly into IAM and policy workflows, while serving an online HTML version of the IPCC reports enriched with confidence scores, timestamps, and provenance. Rather than treating synthesis as episodic and frozen, the project establishes it as an ongoing, dynamic process closing the gap between rapid scientific advances and the urgent needs of global climate policy. This is not merely a tool that can be built, but one that *must* be built if climate science is to remain actionable in real time.

Generalisation.   While initially designed for IPCC assessments, the pipeline generalises to any *seed document* to be expanded or synthesized. By seeding it with a reference report, the system can automatically retrieve, filter, and integrate subsequent literature, maintaining a living version of the document in any evidence intensive domain from biodiversity assessments to public health guidelines.

4 Methods

The workflow will consist of three main stages:

1. **Identification of relevant literature.** We will begin with a seed set of references from the published IPCC reports, including both chapter reference lists and annex bibliographies. Using bibliometric datasets such as OpenAlex and Web of Science, the system will retrieve all forward citations (papers citing the seed set), backward citations (papers cited by the seed set), and semantically similar studies identified via embedding search. Each candidate study will then be evaluated for relevance using a large language model (LLM) relevance classifier, which will process topic-abstract pairs and produce binary inclusion decisions. Relevance will also be checked using cocitation strength with original references.

Only studies receiving a positive classification will proceed to the synthesis stage, ensuring both topical precision and comprehensive retrieval beyond what manual searches achieve.

2. **Synthesis and validation of findings.** For each filtered paper, we will extract key findings relevant to the target IPCC section using LLM based summarisation tuned for factual retention. Extracted statements will be linked to their exact location in the source PDF to preserve provenance. The system will also classify each statement to the correct chapter, subsection, and line of the IPCC text (LOC). A separate, independent AI model differing in architecture and training corpus will then re-parse these passages to verify factual correctness and contextual alignment. Confidence scores will be assigned and periodically recalibrated. Only findings confirmed by this validator will be included in the synthesis. This two-model approach will reduce correlated errors, increase factual accuracy, and provide a transparent chain from claim to source text.

3. **Integration into a live, online IPCC report.** Validated findings will be inserted into a continuously maintained HTML version of the IPCC report, built with a Python backend and a lightweight React-based front-end. Newly integrated studies will be highlighted visually, accompanied by their validation status, date of integration, and direct links to source documents. The report will be query-able by topic, confidence, and publication date, with structured outputs made available via an API for integration into IIASA's Integrated Assessment Models (IAMs) and related policy analysis workflows.

Finally, humans-in-the-loop will check new additions and validate those where AI agents are uncertain. This design will address both timeliness and trust gaps left open by current approaches while substantially reducing the human cost of synthesis. At a conservative €50 per fully loaded research hour, screening labour alone in AR6 likely exceeded €2 million per Working Group. Our hybrid human-AI workflow will aim to cut this cost while enabling more frequent, iteration friendly synthesis cycles without sacrificing rigour. Updates will be comprehensive, auditable, and bias-aware, in contrast to fast but opaque alternatives.

References

1. S. Arrhenius, "On the influence of carbonic acid in the air upon the temperature of the ground," *Philosophical Magazine and Journal of Science*, vol. 41, no. 251, pp. 237–276, 1896.

2. National Research Council, Carbon Dioxide and Climate Committee, *Carbon Dioxide and Climate: A Scientific Assessment*. Washington, D.C.: National Academies Press, 1979.

3. S. Manabe and R. T. Wetherald, "Thermal equilibrium of the atmosphere with a given distribution of relative humidity," *Journal of the Atmospheric Sciences*, vol. 24, no. 3, pp. 241–259, 1967.

4. N. Stern, *The Economics of Climate Change: The Stern Review*. Cambridge University Press, 2006.

5. V. Masson-Delmotte, P. Zhai, H.-O. Pörtner, D. Roberts, J. Skea, P. R. Shukla, A. Pirani, W. Moufouma-Okia, C. Péan, R. Pidcock, S. Connors, J. B. R. Matthews, Y. Chen, X. Zhou, M. I. Gomis, E. Lonnoy, T. Maycock, M. Tignor, and T. Waterfield, eds., *Global Warming of 1.5 °C. An IPCC Special Report on the impacts of global warming of 1.5 °C above pre-industrial levels and related global greenhouse gas emission pathways, in the context of strengthening the global response to the threat of climate change, sustainable development, and efforts to eradicate poverty*. Geneva, Switzerland: Intergovernmental Panel on Climate Change, 2018.

6. J. Rockström, W. Steffen, K. Noone, Persson, F. S. I. Chapin, E. Lambin, T. M. Lenton, M. Scheffer, C. Folke, H. J. Schellnhuber, B. Nykvist, C. A. de Wit, T. Hughes, S. van der Leeuw, H. Rodhe, S. Sörlin, P. K. Snyder, R. Costanza, U. Svedin, M. Falkenmark, L. Karlberg, R. W. Corell, V. J. Fabry, J. Hansen, B. Walker, D. Liverman, K. Richardson, P. Crutzen, and J. Foley, "A safe operating space for humanity," *Nature*, vol. 461, no. 7263, pp. 472–475, 2009.

7. N. Castree, R. Bellamy, and S. Osaka, "The future of global environmental assessments: Making a case for fundamental change," *The Anthropocene Review*, vol. 8, no. 1, pp. 56–82, 2020.

8. J. Jabbour and C. Flachsland, "40 years of global environmental assessments: a retrospective analysis," *Environmental Science & Policy*, vol. 77, pp. 193–202, 2017.

9. T. Ourbak and L. Tubiana, "Changing the game: the paris agreement and the role of scientific communities," *Climate Policy*, vol. 17, no. 7, pp. 819–824, 2017.

10. IPCC, "Summary for policymakers," *Climate Change 2023: Synthesis Report. Contribution of Working Groups I, II and III to the Sixth Assessment Report of the Intergovernmental Panel on Climate Change*, 2023.

11. M. Vardy, M. Oppenheimer, N. K. Dubash, J. O'Reilly, D. Jamieson, *et al.*, "The intergovernmental panel on climate change: Challenges and opportunities," *Annual Review of Environment and Resources*, vol. 42, pp. 55–75, 2017.

12. S. Fortunato, C. T. Bergstrom, K. Börner, J. A. Evans, D. Helbing, S. Milojević, A. M. Petersen, F. Radicchi, R. Sinatra, B. Uzzi, A. Vespignani, D. Wang, and A.-L. Barabási, "Science of science," *Science*, vol. 359, no. 6379, p. eaao0185, 2018.

13. L. Bornmann and R. Mutz, "Growth rates of modern science: A bibliometric analysis based on the number of publications and cited references," *Journal of the Association for Information Science and Technology*, vol. 72, no. 9, pp. 1099–1112, 2021.

14. Intergovernmental Panel on Climate Change, "About the ipcc: How it is structured," 2024. IPCC comprehensive assessment reports are produced roughly every 5–7 years.

15. IPCC, "Climate change 2014: Synthesis report," 2015. AR5 Synthesis Report, Accessed 2025-08-16.

16. IPCC, "Ar6 factsheet 2022." Intergovernmental Panel on Climate Change, 2022. Summary factsheet on publications assessed in AR6. Accessed 2025-08-16.

17. UNFCCC Secretariat, "Summary report on the sbsta–ipcc special event on the working group i contribution to ar6," 2022. AR6 WGI scale: 234 authors; >14,000 publications assessed; >78,000 review comments.

18. UNFCCC Secretariat, "Summary report on the sbsta–ipcc special event: Unpacking the new scientific knowledge and key findings in the working group iii contribution to ar6," 2022. AR6 WGIII scale: 278 lead authors; 354 contributing authors; >18,000 publications; >59,000 comments.

19. D. Khojasteh, M. Haghani, A. Shamsipour, C. Zwack, W. Glamore, *et al.*, "Climate change science is evolving toward adaptation and mitigation solutions," *Wiley Interdisciplinary Reviews: Climate Change*, vol. 15, no. 2, p. e884, 2024.

20. M. Callaghan, J. Minx, and P. Forster, "A topography of climate change research," *Nature Climate Change*, vol. 10, pp. 118–123, 2020.

21. Intergovernmental Panel on Climate Change, "Ipcc trust fund programme and budget (2024 income/expenditure; 2025 proposed)," 2025. Shows 2024 expenditure CHF 5,505,000.

22. M. Callaghan, C.-F. Schleussner, S. Nath, Q. Lejeune, T. R. Knutson, M. Reichstein, G. Hansen, E. Theokritoff, M. Andrijevic, R. J. Brecha, M. Hegarty, C. Jones, K. Lee, A. Lucas, N. van Maanen, I. Menke, P. Pfleiderer, B. Yesil, and J. C. Minx, "Machine-learning-based evidence and attribution mapping of 100,000 climate impact studies," *Nature Climate Change*, vol. 11, no. 11, p. 966–972, 2021.

23. A. J. Sietsma, E. Theokritoff, R. G. Biesbroek, I. Villaverde Canosa, A. Thomas, M. Callaghan, J. C. Minx, and J. D. Ford, "Machine learning evidence map reveals global differences in adaptation action," *One Earth*, vol. 7, no. 2, pp. 280–292, 2024.

24. M. Callaghan, L. Banisch, N. Doebbeling-Hildebrandt, D. Edmondson, C. Flachsland, W. F. Lamb, S. Levi, F. Müller-Hansen, E. Posada, S. Vasudevan, and J. C. Minx, "Machine learning map of climate policy literature reveals disparities between scientific attention, policy density, and emissions," *npj Climate Action*, vol. 4, p. 7, 2025.

25. S. Lück, M. Callaghan, M. Borchers, A. Cowie, S. Fuss, M. Gidden, J. Hartmann, C. Kammann, D. P. Keller, F. Kraxner, W. F. Lamb, N. Mac Dowell, F. Müller-Hansen, G. F. Nemet, B. S. Probst, P. Renforth, T. Repke, W. Rickels, I. Schulte, P. Smith, S. M. Smith, D. Thrän, T. G. Troxler, V. Sick, M. van der Spek, and J. C. Minx, "Scientific literature on carbon dioxide removal revealed as much larger through ai-enhanced systematic mapping," *Nature Communications*, vol. 16, p. 6632, 2025.

26. A. Sharifi, Y. Jin, and P. D. Aboagye, "Mapping and topic modeling of over 133,000 studies on climate change and cities: Trends, geographical disparities, and future research needs," *Urban Climate*, vol. 61, p. 102448, 2025.

27. A.-M. Niculae, S.-V. Oprea, A.-G. Văduva, A. Bâra, and A.-I. Andreescu, "Assessing the role of machine learning in climate research publications," *Sustainability*, vol. 16, no. 24, p. 11086, 2024.

28. N. Webersinke, M. Kraus, J. A. Bingler, and M. Leippold, "Climatebert: A pretrained language model for climate-related text," *arXiv preprint arXiv:2110.12010*, 2021.

29. L. Berrang-Ford, A. R. Siders, A. Lesnikowski, A. P. Fischer, M. W. Callaghan, *et al.*, "A systematic global stocktake of evidence on human adaptation to climate change," *Nature Climate Change*, vol. 11, pp. 989–1000, 2021.

30. L. Berrang-Ford, A. Sietsma, M. Callaghan, J. Minx, P. Scheelbeek, *et al.*, "Systematic mapping of global research on climate and health: a machine learning review," *The Lancet Planetary Health*, vol. 5, no. 8, pp. e514–e525, 2021.

31. A. Thomas, E. Theokritoff, A. Lesnikowski, D. Reckien, K. Jagannathan, *et al.*, "Global evidence of constraints and limits to human adaptation," *Regional Environmental Change*, vol. 21, no. 85, 2021.

32. A. J. Sietsma, R. W. Groenendijk, and R. Biesbroek, "Progress on climate action: a multilingual machine learning analysis of the global stocktake," *Climatic Change*, vol. 176, no. 173, 2023.

33. D. Thulke, Y. Gao, P. Pelser, R. Brune, R. Jalota, *et al.*, "Climategpt: towards ai synthesizing interdisciplinary research on climate change," *arXiv preprint arXiv:2401.09646*, 2024.

34. J. Elliott, R. Lawrence, J. C. Minx, O. T. Oladapo, P. Ravaud, *et al.*, "Decision makers need 'living' evidence synthesis," *Nature*, vol. 600, no. 7889, pp. 383–385, 2021.

35. R. van de Schoot, J. de Bruin, R. Schram, P. Zahedi, J. de Boer, *et al.*, "An open source machine learning framework for efficient and transparent systematic reviews," *Nature Machine Intelligence*, vol. 3, pp. 125–133, 2021.

36. M. D. Mastrandrea, K. J. Mach, G.-K. Plattner, O. Edenhofer, T. F. Stocker, *et al.*, "The ipcc ar5 guidance note on consistent treatment of uncertainties: a common approach across the working groups," *Climatic Change*, vol. 108, no. 4, pp. 675–691, 2011.

37. M. Leippold, S. Vaghefi, D. Stammbach, V. Muccione, J. A. Bingler, *et al.*, "Automated fact-checking of climate change claims with large language models," *arXiv preprint arXiv:2401.12566*, 2024.

38. M. Majszak and J. Jebeile, "Expert judgment in climate science: how it is used and how it can be justified," *Studies in History and Philosophy of Science*, vol. 100, pp. 32–38, 2023.

39. K. Wu *et al.*, "An automated framework for assessing how well llms cite relevant medical references," *Nature Communications*, vol. 16, no. ???, p. ???, 2025. Finds 30–90% of statements unsupported across popular LLMs.

40. L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin, and T. Liu, "A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions," *ACM Transactions on Information Systems*, 2024.

41. J. Clark and colleagues, "Generative artificial intelligence use in evidence synthesis: a systematic review," *Research Synthesis Methods*, pp. 601–619, 2025.

42. M. Stede and R. Patz, "The climate change debate and natural language processing," in *Proceedings of the 1st Workshop on NLP for Positive Impact*, pp. 8–18, 2021.

43. "Semantic scholar." https://www.semanticscholar.org.

44. "Elicit." https://elicit.com.

45. "Consensus." https://consensus.app.

46. "Paperguide." https://paperguide.ai.

47. "Scite." https://scite.ai.

48. "Connected papers." https://connectedpapers.com.

49. "Researchrabbit." https://researchrabbit.ai.

50. "Scholarcy." https://scholarcy.com.

51. "Scispace." https://scispace.com.

52. "Iris.ai." https://iris.ai.

53. IPCC Working Group III, "Technical summary. climate change 2022: Mitigation of climate change," 2022.

54. IPCC Working Group III, "Carbon dioxide removal (cdr) factsheet," 2022.

55. IPCC Working Group III, "Chapter 5: Demand, services and social aspects of mitigation," in *Climate Change 2022: Mitigation of Climate Change*, 2022.

56. IPCC Working Group III, "Annex iii: Scenarios and modelling methods," 2022.

57. G. P. Peters and colleagues, "Ar6 scenarios database: an assessment of current curation and use," *npj Climate Action*, 2023.

58. B. Cointe and colleagues, "The ar6 scenario explorer and the history of ipcc scenario databases," *npj Climate Action*, 2024.

59. IIASA, "Ar6 scenario explorer and database," 2022.

60. IPCC Working Group II, "Summary for policymakers. climate change 2022: Impacts, adaptation and vulnerability," 2022.

61. IPCC Working Group II, "Chapter 6: Cities, settlements and key infrastructure," 2022.

62. IPCC Working Group I, "Chapter 11: Weather and climate extreme events in a changing climate," in *Climate Change 2021: The Physical Science Basis*, 2021.

63. B. Uzzi, S. Mukherjee, M. Stringer, and B. Jones, "Atypical combinations and scientific impact," *Science*, vol. 342, no. 6157, pp. 468–472, 2013.